

ML Day2

March 18, 2023

1 Handling Missing values

```
[1]: import seaborn as sns
```

```
[2]: df=sns.load_dataset("titanic")
```

```
[3]: df
```

```
[3]:      survived  pclass    sex  age  sibsp  parch   fare embarked  class \
0            0        3   male  22.0     1     0   7.2500         S   Third
1            1        1  female  38.0     1     0  71.2833         C   First
2            1        3  female  26.0     0     0   7.9250         S   Third
3            1        1  female  35.0     1     0  53.1000         S   First
4            0        3   male  35.0     0     0   8.0500         S   Third
..          ...      ...    ...  ...  ...    ...    ...    ...    ...
886           0        2   male  27.0     0     0  13.0000         S  Second
887           1        1  female  19.0     0     0  30.0000         S   First
888           0        3  female   NaN     1     2  23.4500         S   Third
889           1        1   male  26.0     0     0  30.0000         C   First
890           0        3   male  32.0     0     0   7.7500         Q   Third
```

```
      who  adult_male deck  embark_town  alive  alone
0     man         True  NaN  Southampton    no  False
1  woman        False   C   Cherbourg   yes  False
2  woman        False  NaN  Southampton   yes   True
3  woman        False   C   Southampton   yes  False
4     man         True  NaN  Southampton    no   True
..    ...      ...    ...    ...    ...    ...
886   man         True  NaN  Southampton    no   True
887  woman        False   B  Southampton   yes   True
888  woman        False  NaN  Southampton    no  False
889   man         True   C   Cherbourg   yes   True
890   man         True  NaN  Queenstown    no   True
```

[891 rows x 15 columns]

```
[4]: df.isna()
```

```
[4]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	\
0	False	False	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	False	False	
..	
886	False	False	False	False	False	False	False	False	False	
887	False	False	False	False	False	False	False	False	False	
888	False	False	False	True	False	False	False	False	False	
889	False	False	False	False	False	False	False	False	False	
890	False	False	False	False	False	False	False	False	False	

	who	adult_male	deck	embark_town	alive	alone
0	False	False	True	False	False	False
1	False	False	False	False	False	False
2	False	False	True	False	False	False
3	False	False	False	False	False	False
4	False	False	True	False	False	False
..
886	False	False	True	False	False	False
887	False	False	False	False	False	False
888	False	False	True	False	False	False
889	False	False	False	False	False	False
890	False	False	True	False	False	False

[891 rows x 15 columns]

```
[5]: df.isna().sum()
```

```
[5]: survived      0
pclass            0
sex              0
age             177
sibsp           0
parch           0
fare            0
embarked         2
class            0
who             0
adult_male       0
deck           688
embark_town      2
alive           0
alone           0
dtype: int64
```

```
[6]: df.drop(["deck"],axis=1,inplace=True)
```

```
[7]: df
```

```
[7]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class \
0	0	3	male	22.0	1	0	7.2500	S	Third
1	1	1	female	38.0	1	0	71.2833	C	First
2	1	3	female	26.0	0	0	7.9250	S	Third
3	1	1	female	35.0	1	0	53.1000	S	First
4	0	3	male	35.0	0	0	8.0500	S	Third
..
886	0	2	male	27.0	0	0	13.0000	S	Second
887	1	1	female	19.0	0	0	30.0000	S	First
888	0	3	female	NaN	1	2	23.4500	S	Third
889	1	1	male	26.0	0	0	30.0000	C	First
890	0	3	male	32.0	0	0	7.7500	Q	Third

	who	adult_male	embark_town	alive	alone
0	man	True	Southampton	no	False
1	woman	False	Cherbourg	yes	False
2	woman	False	Southampton	yes	True
3	woman	False	Southampton	yes	False
4	man	True	Southampton	no	True
..
886	man	True	Southampton	no	True
887	woman	False	Southampton	yes	True
888	woman	False	Southampton	no	False
889	man	True	Cherbourg	yes	True
890	man	True	Queenstown	no	True

[891 rows x 14 columns]

```
[8]: df.isna().sum()
```


```
[8]:
```

survived	0
pclass	0
sex	0
age	177
sibsp	0
parch	0
fare	0
embarked	2
class	0
who	0
adult_male	0
embark_town	2
alive	0

```
alone          0
dtype: int64
```

```
[9]: sns.histplot(df["age"],kde=True)
```

```
[9]: <AxesSubplot: xlabel='age', ylabel='Count'>
```



```
[10]: df["age"].fillna(df["age"].mean(),inplace=True)
```

```
[11]: df
```

```
[11]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	\
0	0	3	male	22.000000	1	0	7.2500	S	
1	1	1	female	38.000000	1	0	71.2833	C	
2	1	3	female	26.000000	0	0	7.9250	S	
3	1	1	female	35.000000	1	0	53.1000	S	
4	0	3	male	35.000000	0	0	8.0500	S	
..		
886	0	2	male	27.000000	0	0	13.0000	S	
887	1	1	female	19.000000	0	0	30.0000	S	
888	0	3	female	29.699118	1	2	23.4500	S	

889	1	1	male	26.000000	0	0	30.0000	C
890	0	3	male	32.000000	0	0	7.7500	Q

	class	who	adult_male	embark_town	alive	alone
0	Third	man	True	Southampton	no	False
1	First	woman	False	Cherbourg	yes	False
2	Third	woman	False	Southampton	yes	True
3	First	woman	False	Southampton	yes	False
4	Third	man	True	Southampton	no	True
..
886	Second	man	True	Southampton	no	True
887	First	woman	False	Southampton	yes	True
888	Third	woman	False	Southampton	no	False
889	First	man	True	Cherbourg	yes	True
890	Third	man	True	Queenstown	no	True

[891 rows x 14 columns]

```
[12]: df.isna().sum()
```

```
[12]: survived      0
pclass            0
sex              0
age              0
sibsp            0
parch            0
fare             0
embarked         2
class            0
who              0
adult_male       0
embark_town      2
alive            0
alone            0
dtype: int64
```

```
[13]: median=df["embarked"].notna().mode()[0]
```

```
[14]: df["embarked"].fillna(median,inplace=True)
```

```
[15]: df.isna().sum()
```

```
[15]: survived      0
pclass            0
sex              0
age              0
sibsp            0
```

```

parch      0
fare       0
embarked   0
class      0
who        0
adult_male 0
embark_town 2
alive      0
alone      0
dtype: int64

```

```
[16]: df.embark_town.fillna("Missing",inplace=True)
```

```
[17]: df.isna().sum()
```

```

[17]: survived      0
pclass            0
sex              0
age              0
sibsp            0
parch            0
fare             0
embarked         0
class            0
who              0
adult_male       0
embark_town      0
alive           0
alone           0
dtype: int64

```

```
[18]: df
```

```

[18]:   survived  pclass  sex      age  sibsp  parch  fare  embarked  \
0         0      3  male  22.000000      1      0   7.2500         S
1         1      1 female  38.000000      1      0  71.2833         C
2         1      3 female  26.000000      0      0   7.9250         S
3         1      1 female  35.000000      1      0  53.1000         S
4         0      3  male  35.000000      0      0   8.0500         S
..      ...    ...   ...    ...    ...    ...   ...
886        0      2  male  27.000000      0      0  13.0000         S
887        1      1 female  19.000000      0      0  30.0000         S
888        0      3 female  29.699118      1      2  23.4500         S
889        1      1  male  26.000000      0      0  30.0000         C
890        0      3  male  32.000000      0      0   7.7500         Q

      class  who  adult_male  embark_town  alive  alone

```

0	Third	man	True	Southampton	no	False
1	First	woman	False	Cherbourg	yes	False
2	Third	woman	False	Southampton	yes	True
3	First	woman	False	Southampton	yes	False
4	Third	man	True	Southampton	no	True
..
886	Second	man	True	Southampton	no	True
887	First	woman	False	Southampton	yes	True
888	Third	woman	False	Southampton	no	False
889	First	man	True	Cherbourg	yes	True
890	Third	man	True	Queenstown	no	True

[891 rows x 14 columns]

2 Handling imbalanced Data

2.1 Making Data

```
[19]: import numpy as np
import pandas as pd

# Set the random seed for reproducibility
np.random.seed(123)

# Create a dataframe with two classes
n_samples = 1000
class_0_ratio = 0.9
n_class_0 = int(n_samples * class_0_ratio)
n_class_1 = n_samples - n_class_0

## CREATE MY DATAFRAME WITH IMBALANCED DATASET
class_0 = pd.DataFrame({
    'feature_1': np.random.normal(loc=0, scale=1, size=n_class_0),
    'feature_2': np.random.normal(loc=0, scale=1, size=n_class_0),
    'target': [0] * n_class_0
})

class_1 = pd.DataFrame({
    'feature_1': np.random.normal(loc=2, scale=1, size=n_class_1),
    'feature_2': np.random.normal(loc=2, scale=1, size=n_class_1),
    'target': [1] * n_class_1
})

df=pd.concat([class_0,class_1]).reset_index(drop=True)
```

```
[20]: df
```

```
[20]:
```

	feature_1	feature_2	target
0	-1.085631	0.551302	0
1	0.997345	0.419589	0
2	0.282978	1.815652	0
3	-1.506295	-0.252750	0
4	-0.578600	-0.292004	0
..
995	1.376371	2.845701	1
996	2.239810	0.880077	1
997	1.131760	1.640703	1
998	2.902006	0.390305	1
999	2.697490	2.013570	1

[1000 rows x 3 columns]

```
[21]: df["target"].value_counts()
```

```
[21]: 0    900
      1    100
      Name: target, dtype: int64
```

```
[22]: major=df[df["target"]==0]
      minor=df[df["target"]==1]
```

```
[23]: major.shape,minor.shape
```

```
[23]: ((900, 3), (100, 3))
```

```
[24]: from sklearn.utils import resample
      minor2=resample(minor,replace=True,n_samples=len(major),random_state=10)
```

```
[25]: minor2
```

```
[25]:
```

	feature_1	feature_2	target
909	3.239635	1.361938	1
915	3.519471	-0.233905	1
964	2.397060	0.740228	1
928	1.868135	1.026563	1
989	3.013493	2.047240	1
..
936	3.727988	3.468919	1
928	1.868135	1.026563	1
947	1.402209	2.775845	1
919	1.804892	2.842652	1
902	1.795683	1.803557	1

[900 rows x 3 columns]


```
[26]: minor2.shape
```

```
[26]: (900, 3)
```

```
[27]: major2=resample(major,replace=False,n_samples=len(minor),random_state=10)
```

```
[28]: major2
```

```
[28]:
```

	feature_1	feature_2	target
437	-1.639397	0.273073	0
131	-1.100043	1.191189	0
633	0.600571	0.627744	0
195	-3.231055	-1.725890	0
230	-1.600441	-0.304086	0
..
235	-0.434167	-0.265576	0
192	0.199582	-0.096391	0
775	0.048109	-0.805562	0
718	0.301290	0.907483	0
769	-1.094273	0.639969	0

```
[100 rows x 3 columns]
```

```
[29]: minor.shape,major2.shape
```

```
[29]: ((100, 3), (100, 3))
```

```
[30]: major.shape,minor2.shape
```

```
[30]: ((900, 3), (900, 3))
```

3 SMOTE

```
[31]: !pip install imblearn
```

```
Requirement already satisfied: imblearn in /opt/conda/lib/python3.10/site-  
packages (0.0)  
Requirement already satisfied: imbalanced-learn in  
/opt/conda/lib/python3.10/site-packages (from imblearn) (0.10.1)  
Requirement already satisfied: scipy>=1.3.2 in /opt/conda/lib/python3.10/site-  
packages (from imbalanced-learn->imblearn) (1.9.3)  
Requirement already satisfied: threadpoolctl>=2.0.0 in  
/opt/conda/lib/python3.10/site-packages (from imbalanced-learn->imblearn)  
(3.1.0)  
Requirement already satisfied: scikit-learn>=1.0.2 in  
/opt/conda/lib/python3.10/site-packages (from imbalanced-learn->imblearn)  
(1.2.0)
```

Requirement already satisfied: joblib>=1.1.1 in /opt/conda/lib/python3.10/site-packages (from imbalanced-learn->imblearn) (1.2.0)
Requirement already satisfied: numpy>=1.17.3 in /opt/conda/lib/python3.10/site-packages (from imbalanced-learn->imblearn) (1.23.5)

```
[32]: from imblearn.over_sampling import SMOTE
```

```
[33]: ans=SMOTE()
```

```
[34]: df["target"].value_counts()
```

```
[34]: 0    900  
      1    100  
      Name: target, dtype: int64
```

```
[35]: x,y=ans.fit_resample(df[["feature_1","feature_2"]],df["target"])
```

```
[36]: y
```

```
[36]: 0      0  
      1      0  
      2      0  
      3      0  
      4      0  
      ..  
     1795     1  
     1796     1  
     1797     1  
     1798     1  
     1799     1  
      Name: target, Length: 1800, dtype: int64
```

```
[37]: x,y
```

```
[37]: (      feature_1  feature_2  
      0    -1.085631    0.551302  
      1     0.997345    0.419589  
      2     0.282978    1.815652  
      3    -1.506295   -0.252750  
      4    -0.578600   -0.292004  
      ...      ...      ...  
     1795    1.090655    2.299404  
     1796    1.894705    3.051485  
     1797    2.391335    3.093234  
     1798    2.392714    0.742123  
     1799    2.276174    2.566088
```

```
[1800 rows x 2 columns],
0      0
1      0
2      0
3      0
4      0
..
1795    1
1796    1
1797    1
1798    1
1799    1
Name: target, Length: 1800, dtype: int64)
```

```
[38]: import pandas as pd
```

```
[39]: pd.DataFrame(x)
pd.DataFrame(y)
df=pd.concat([x,y],axis=1)
```

```
[40]: df
```

```
[40]:
```

	feature_1	feature_2	target
0	-1.085631	0.551302	0
1	0.997345	0.419589	0
2	0.282978	1.815652	0
3	-1.506295	-0.252750	0
4	-0.578600	-0.292004	0
...
1795	1.090655	2.299404	1
1796	1.894705	3.051485	1
1797	2.391335	3.093234	1
1798	2.392714	0.742123	1
1799	2.276174	2.566088	1

```
[1800 rows x 3 columns]
```

```
[41]: df["target"].value_counts()
```

```
[41]: 0    900
1    900
Name: target, dtype: int64
```

4 Interpolation


4.1 Linear Interpolate

```
[42]: x=np.array([1,2,3,4,5])  
      y=np.array([1,2,3,4,5])
```

```
[43]: import matplotlib.pyplot as plt
```

```
[44]: plt.scatter(x,y)
```

```
[44]: <matplotlib.collections.PathCollection at 0x7f6c8aa962c0>
```



```
[45]: x_new=np.linspace(1,5,10)
```

```
[46]: x_new
```

```
[46]: array([1.          , 1.44444444, 1.88888889, 2.33333333, 2.77777778,  
        3.22222222, 3.66666667, 4.11111111, 4.55555556, 5.          ])
```


```
[47]: y_new=np.interp(x_new,x,y)
```

```
[48]: y_new
```

```
[48]: array([1.          , 1.44444444, 1.88888889, 2.33333333, 2.77777778,  
          3.22222222, 3.66666667, 4.11111111, 4.55555556, 5.          ])
```

```
[49]: plt.scatter(x_new,y_new)
```

```
[49]: <matplotlib.collections.PathCollection at 0x7f6c8278f340>
```




4.2 Cubic Interpolate

```
[50]: a=np.array([1,2,3,4,5])  
      b=np.array([2,5,7,11,23])
```

```
[51]: plt.scatter(a,b)
```

```
[51]: <matplotlib.collections.PathCollection at 0x7f6c82617910>
```



```
[52]: from scipy.interpolate import interp1d
```


```
[53]: f=interp1d(a,b,kind="cubic")
```

```
[54]: a_new=np.linspace(1,5,35)
```

```
[55]: b_new=f(a_new)
```

```
[56]: plt.scatter(a_new,b_new)
```

```
[56]: <matplotlib.collections.PathCollection at 0x7f6c826aa050>
```



5 Percentiles and Outliers

```
[57]: marks=[11,23,56,80,90,90,100,100,100,109,112,132,148,179,230]
```

```
[58]: minimum,q1,median,q3,maximum=np.quantile(marks,[0,0.25,0.50,0.75,1])
```

```
[59]: minimum,q1,median,q3,maximum
```

```
[59]: (11.0, 85.0, 100.0, 122.0, 230.0)
```

```
[60]: IQR=q3-q1
IQR
```

```
[60]: 37.0
```


```
[61]: lower_fence=IQR-(1.5*q1)
higher_fence=IQR+(1.5*q3)
```

```
[62]: lower_fence,higher_fence
```

```
[62]: (-90.5, 220.0)
```

```
[63]: sns.boxplot(marks)
```

```
[63]: <AxesSubplot: >
```



```
[ ]:
```